

ADRIÁN SALVADOR EKOMO MESI OBONO

Leipzig, Saxony, Germany · adriansalvadorekomo@outlook.com · [linkedin.com/in/adriansalvadorekomo](https://www.linkedin.com/in/adriansalvadorekomo)
github.com/adriansalvadorekomo · adriansalvadorekomodev.vercel.app · **Available from December 2026**

Engineering student building the platforms that data and ML workloads run on: reproducible infrastructure, quality-gated CI/CD and governed data pipelines. I eliminate manual, error-prone steps where infrastructure, CI/CD and data intersect. Seeking a **Data Platform / Data Engineering / DevOps internship**.

TECHNICAL SKILLS

Platform & Infrastructure

Infrastructure as Code (Terraform, Ansible), OpenStack private cloud, Linux servers, Nginx reverse proxy & HTTPS

Delivery & Automation

CI/CD with quality gates, Jenkins, SonarQube, Git-based review workflows

Containers & Orchestration

Docker, Docker Compose, Kubernetes (minikube, kubectl)

Data Engineering

Apache Spark (PySpark, Spark SQL, Spark Streaming), Databricks, medallion lakehouse, data-quality gates, star-schema warehousing

Pipelines & Observability

Apache Airflow, Prefect, n8n, MLflow, Prometheus, Grafana

Programming

Python, SQL, Bash, Java

PROFESSIONAL EXPERIENCE

RIF – Rassemblement des Ingénieurs Francophones | DevOps Intern · France (Remote) · Jul 2026 – Sep 2026

Operated cloud infrastructure for a volunteer organization of 10+ engineers, replacing hand-built servers with reproducible code.

- Replaced manual provisioning with code: defined OpenStack infrastructure (VMs, networks, Floating IPs, Security Groups) in Terraform and server setup as idempotent Ansible roles (Docker, UFW, reverse proxy).
- Shipped a containerized Open edX platform behind Nginx with HTTPS and DuckDNS, making the deployment repeatable.
- Designed a PRIMARY→BACKUP bastion setup that restores SSH access on a single Floating IP in ~5 s during failover tests.
- Kept every infrastructure change reviewable through Git-based version control.

PROJECTS

DataOps Lakehouse & Event Analytics | Solo + team builds · Jun 2026 – Aug 2026

A data platform provisioned as code, with a data-quality gate before the Gold layer.

- Solo: built a PostgreSQL + Databricks medallion lakehouse with PySpark transformations (Bronze → Silver → DQ gate → Gold) under Terraform with PR-gated CI; end-to-end run green, reconciling 99.4M of synthetic sample revenue. [\[code\]](#)
- In a 6-person academic team: delivered a star-schema warehouse unifying 3 source systems, Prophet demand forecasting (MAPE under 15%) and 11 Dockerized microservices behind nginx, with MLflow tracking and Prometheus + Grafana monitoring; orchestrated pipelines with Apache Airflow and automated ML training with n8n. [\[code\]](#)
- Solo: Spark, Spark Streaming, PySpark and Spark SQL notebooks, public on [GitHub](#).

Project Management App | Solo build · Mar 2026 – Jun 2026

Made a 3-tier app (Spring Boot, Angular, MySQL) reproducible with one command from commit to running stack, through a 10-stage CI/CD pipeline with quality gates, then ported it to Kubernetes. [\[code\]](#)

- Eliminated “works on my machine” failures with a 5-service Compose stack: the backend starts only after a healthy MySQL, secrets travel via environment and the database exposes no public port.
- Built a 10-stage Jenkins pipeline that fails fast: build, SonarQube analysis, Quality Gate (red blocks the release), multistage Docker builds, immutable :build-N tags, registry pull-back proof and redeploy verified with curl checks.
- Ported the stack to Kubernetes on a local minikube cluster, managed with kubectl and declarative manifests (academic exercise).
- Kept quality infrastructure sound: dedicated SonarQube stack on PostgreSQL (9.9+ dropped MySQL) and credentials that never touch code (Jenkins secrets, agent-level registry auth).

Customer Churn MLOps | Solo build · Nov 2025 – Jan 2026

Turned a one-off churn notebook into a tested, scheduled MLOps pipeline that retrains and re-scores every night. [\[code\]](#)

- Stopped bad code from ever training a model: separated pure pipeline logic from Prefect orchestration and blocked runs with ruff, bandit and 6 pytest gates.
- Automated retraining and re-scoring without overlap using three staggered Prefect flows (full run 02:00, retrain 04:00, rescore 06:00 UTC).
- Reached 85.5% accuracy (RandomForest, 2,000-customer holdout) and documented the weak churner recall (0.42) with class imbalance as the next fix, instead of hiding it.
- Shipped a versioned joblib model with a one-command run (`uv run python main.py`).

EDUCATION

ESPRIT School of Engineering | Engineering – Computer Science, ERP-BI / Data Engineering (5th year) · Sep 2022 – Dec 2026

Philipps-Universität Marburg | Exchange student (Freemover), Dept. of Mathematics and Computer Science · Oct 2026 – present

CERTIFICATIONS & LANGUAGES

Certifications: Get Started with Databricks for Data Engineering, Databricks (Sep 2026) · DataCamp Data Engineer Associate, SQL Associate & Python Data Associate (2026)

Languages: Spanish (native) · English (C1) · French (C1) · German (A1, learning)