

Adrian Salvador Ekomo

Data Engineering Student | DevOps → DataOps | **Databricks** · **Terraform** · PySpark · **SQL**

📍 Tunis, Tunisia • ✉ adriansalvadorekomo@outlook.com • 🗣 adriansalvadorekomo • 🌐 Adrian Salvador Ekomo

Professional Summary

Computer Engineering student focused on Data Engineering with a systems-thinking approach — how data moves and where it breaks. DevOps intern (**OpenStack**, **Terraform**, **Ansible**, Open edX) + solo DataOps build (**PostgreSQL** + **Databricks** medallion, DQ gates, **Terraform**, CI). Seeking Junior **Data Engineer** role.

Education

Esprit School of Engineering — **Sept 2022 – June 2027**, Engineering in Computer Science – Tunis, Tunisia

- Specialized in Data Engineering and Business Intelligence

Experience

DevOps Engineer · Internship (Remote), RIF — Rassemblement des Ingénieurs Francophones – La Marsa, Tunis July 2026 – Sept 2026

My scope: **Terraform**/**OpenStack** provisioning, **Ansible** configuration, containerized Open edX deployment, and HA Bastion validation (code).

- Provisioned VMs, networks, Floating IPs and Security Groups with **Terraform** configured **Docker**, UFW and reverse proxy via idempotent **Ansible** roles
- Deployed containerized Open edX with Nginx, DuckDNS and HTTPS, and validated PRIMARY→BACKUP HA Bastion failover in ~5s on a single Floating IP
- Managed secure SSH access and Git-based version control across the deployment pipeline

Projects

Event Analytics Platform

Jan 2026 – May 2026

Academic team project (6 people); my scope: integration, **Docker**/Nginx delivery, and chatbot/API fixes. End-to-end event analytics platform with star-schema warehouse, **ETL**, ML forecasting/anomaly detection, NL-to-**SQL**, and containerized microservices.

- Built a star-schema warehouse in **PostgreSQL** that unified event, booking, and marketing data across 3 source systems for cross-functional analysis — organized around business questions, not source schemas. Used Talend for **ETL** transformations and **Apache Airflow** for reliable ingestion
- Developed ML models — Prophet for demand forecasting (handles event seasonality naturally) with MAPE under 15%, statsmodels for anomaly detection, and scikit-learn for classification — tracked through **MLflow** and exposed via a natural-language interface so anyone could explore the data
- Containerized 11 microservices with **Docker** behind an nginx proxy, fixed chatbot/API integration, and automated monitoring plus ML retraining
- [View code on GitHub](#)

Environmental Monitoring AI

Mar 2026 – Mar 2026

AI-powered environmental monitoring platform developed during a national innovation hackathon, integrating satellite image segmentation, time-series forecasting, air quality prediction, and a multilingual chatbot interface.

- Built satellite segmentation models with **PyTorch** on multispectral imagery to classify soil contamination and detect water turbidity — ~85% accuracy on soil classification, giving field teams real-time environmental data they could act on
- Used **Airbyte** to sync sensor data from 20+ field devices into a central database — eliminating manual data wrangling — then built time-series models with scikit-learn to forecast contamination trends and air quality with RMSE under 20%, turning noisy measurements into clear environmental insights
- [View code on GitHub](#)

DataOps Lakehouse for E-commerce Analytics

Apr 2026 – present

Solo personal DataOps build (e-commerce marketplace): **PostgreSQL** OLTP + **Databricks** Lakehouse (Bronze → Silver → DQ gate → Gold), IaC with **Terraform**.

- Built medallion Lakehouse on **Databricks** Free Edition (Delta, **Unity Catalog**, PySpark/**SQL**): Bronze COPY INTO, 6 Silver entities, fail-fast DQ gate R1–R9, Gold marts (fact_sales, sales_daily, customer_360, inventory_kpis) — orchestrated with Workflows, infra as code in **Terraform**
- Validated end-to-end on a live sample (SUCCESS, 4/4 tasks green): DQ gate R1–R9 passed, 99.4M sample revenue reconciled to full baseline; Git-versioned, PR-gated CI (lakehouse tests, **SQL** guards, terraform validate)
- [View code on GitHub](#)

Skills

Data Engineering: **SQL**, **PostgreSQL**, **Databricks**, **Delta Lake**, **Unity Catalog**, PySpark, **Databricks** Workflows, Talend, **ETL/ELT**, Data Warehousing, Data Modeling, Dimensional Modeling (familiar: **Airflow**, **Airbyte**, **dbt**)

Cloud & DevOps: **Terraform**, **OpenStack**, **Ansible**, **Docker**, **Linux**, Git, GitHub Actions, **FastAPI**, **REST** APIs, Nginx, n8n, Agile

Analytics & BI Engineering: **Databricks SQL**, **Power BI**, Metabase, SAP BW (learning), Prometheus, Grafana

ML/DL Model Integration & MLOps: Scikit-learn, **XGBoost**, statsmodels, **PyTorch**, **MLflow**, NLP-to-**SQL**, Time-Series Forecasting, Clustering, Anomaly Detection

Programming Languages: **Python**, Java, Bash

Languages

Spanish: Native — Full professional proficiency

English: B2 — Read ML/AI papers, write technical docs

French: B2 — Read technical papers, write documentation

German: A1 — Basic

Certifications

Data Engineer Associate — DataCamp (Sept 2026)

Get Started with **Databricks** for Data Engineering — **Databricks** (Sept 2026)

SQL Associate — DataCamp (Sept 2026)

Python Data Associate — DataCamp (Sept 2026)

Prompt Engineering for Everyone — Cognitive Class / IBM (July 2026)